



PyTorch

Researcher's

What is PyTorch ?



- PyTorch is an open-source machine learning framework, primarily used for deep learning
- It's developed by Meta (Facebook) and a large community of contributors
- Think of it as a toolbox filled with specialized tools that make building, training, and deploying machine learning models much easier

Features



Tensor Computation

↪ like numpy

- At its core, PyTorch is built around **tensors** which are multi-dimensional arrays
- **GPU Acceleration**: The major advantage of PyTorch (and other deep learning frameworks) is its ability to leverage GPUs (Graphics Processing Units). GPUs are designed for parallel processing, making them incredibly fast at the matrix operations that are fundamental to deep learning.
- PyTorch seamlessly moves tensors between your CPU and GPU, allowing you to take advantage of this speedup.

Dynamic Computational Graph

- The graph representing your model is built as you execute the code. Each operation (e.g., addition, multiplication, activation function) is added to the graph as it's performed
- **Benefits of Dynamic Graphs**:
 - **Easier Debugging**: Easy to step through code and see exactly what's happening at each stage.
 - **Flexibility**: Dynamic graphs make it easier to create models with variable-length inputs (e.g., natural language processing tasks like machine translation) and complex control flow.
 - **Intuitive**: It feels more natural to write code that reflects how you think about the model

Features



- **Autograd (Automatic Differentiation)**

- Gradients are essential for training neural networks using techniques like backpropagation
- You define your forward pass and autograd handles the backward pass

- **Modules (nn package)**

- Provides a collection of pre-built neural network layers, loss functions and optimization algorithms

- **Data Loading & Preprocessing**

- torchvision: Image processing (datasets, transforms)
- torchtext: Text data (tokenization, vocabulary building)
- torchaudio: Audio processing

- **Distributed Training**

- PyTorch has excellent support for distributed training, allowing you to train models across multiple GPUs or even multiple machines. This is crucial for large-scale projects

Why PyTorch ?



- **Ease of Use & Debugging**: The dynamic graph and Pythonic API make it easier to learn, experiment with, and debug models.
- **Research-Friendly**: PyTorch is widely used in the research community due to its flexibility and ease of modification. Researchers often publish their code in PyTorch, making it accessible to others.
- **Growing Community**: A large and active community provides support, tutorials, and pre-trained models.
- **Production Readiness**: While initially favored for research, PyTorch has matured significantly and is now used in production environments. Tools like TorchScript allow you to optimize and deploy PyTorch models efficiently.
- **Integration with Other Tools**: PyTorch integrates well with other Python libraries and tools, such as NumPy, SciPy, and Jupyter notebooks.

Basic Workflow



- **Define the Model:** Create a neural network using `torch.nn` modules (layers).
- **Prepare Data:** Load and preprocess your data using `torchvision`, `torchtext`, or custom code.
- **Define Loss Function and Optimizer:** Choose a loss function (e.g., `torch.nn.CrossEntropyLoss`) and an optimization algorithm (e.g., `torch.optim.Adam`).
- **Training Loop:** → for multiple Epochs
 - Iterate over your data in batches.
 - Perform a forward pass through the model to get predictions.
 - Calculate the loss (difference between predictions and actual values).
 - Perform a backward pass to calculate gradients.
 - Update the model's parameters using the optimizer.
- **Evaluation:** Evaluate your trained model on a separate dataset to assess its performance.



Tensor

What is Tensor?

■ Multi-Dimensional Array

- At its heart, a tensor is just a multi-dimensional array of numerical data
- Think of it as an extension of NumPy arrays, but with crucial differences related to GPU acceleration and automatic differentiation

→ similar values

■ Generalization of Scalars, Vectors, Matrices

■ Scalar (0-D Tensor)

→ scalar_tensor = ...[5]

- A single number. Example: `torch.tensor(5)`

5

■ Vector (1-D Tensor)

→ 1D array of similar values

→ [1, 2, 3, 4]

- A single row or column of numbers. Example: `torch.tensor([1, 2, 3])`



■ Matrix (2-D Tensor)

- A table of numbers. Example: `torch.tensor([[1, 2], [3, 4]])`

$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$

⇒ $\begin{bmatrix} [1, 2] \\ [3, 4] \end{bmatrix}$



■ Higher-Order Tensors

- Tensors with more than two dimensions
- Think of them as cubes, hypercubes, or even more abstract structures



Creating Tensors



- **torch.tensor()**: The most general way to create a tensor from Python data
- **torch.zeros()**: Creates a tensor filled with zeros
- **torch.ones()**: Creates a tensor filled with ones
- **torch.rand()**: Creates a tensor filled with random numbers from a uniform distribution (between 0 and 1)
- **torch.randn()**: Creates a tensor filled with random numbers from a standard normal distribution (mean 0, variance 1)
- **torch.empty()**: Creates an uninitialized tensor (filled with garbage values). Use this when you're going to immediately fill the tensor with your own data
- **torch.arange()**: Creates a 1-D tensor with values in a specified range
- **torch.linspace()**: Creates a 1-D tensor with evenly spaced values over a specified interval



Key properties of Tensor

- shape: A tuple representing the dimensions of the tensor (rows x cols)
- dtype: The data type of the elements in the tensor int32, int16, int8 ..
- device: Indicates where the tensor is stored (CPU or GPU)
- item(): get data stored in a tensor

size() = No of items present in the tensor

nbytes() = returns total memory required to load tensor

Tensor Operations



- **Broadcasting:** PyTorch automatically handles broadcasting when performing operations on tensors with different shapes, as long as certain conditions are met
- **Arithmetic Operators:** Perform element-wise arithmetic operations
- **Matrix Multiplication (@ or torch.matmul()):** Performs matrix multiplication
- **Reshaping (reshape() or view()):** Changes the shape of a tensor without changing its data
- **Indexing and Slicing:** Access elements or sub-tensors using indexing and slicing
- **Common Functions:** PyTorch provides a wide range of functions for tensor manipulation
 - **torch.max(), torch.min():** Find the maximum and minimum values
 - **torch.mean(), torch.sum():** Calculate the mean and sum of elements
 - **torch.exp(), torch.log():** Exponential and logarithm functions
 - **torch.transpose():** Transpose a tensor (swap rows and columns)

$$t = [1, 2, 3, 4, 5]$$

broadcasting

$$* \text{ addition} = t + 10 = [1+10, 2+10, 3+10, 4+10] = [11, 12, 13, 14]$$

$$* \text{ subtraction} = t - 5 = [-4, -3, -2, -1, 0]$$

$$* \text{ power} = t ** 2 = [1, 4, 9, 16, 25]$$



Important Considerations

■ GPU Acceleration

- Move tensors to the GPU using `.to("cuda")` (if a CUDA-enabled GPU is available) for significant performance gains in deep learning computations

■ In-Place Operations

- Some PyTorch functions have an out= argument that allows you to modify the tensor in-place (without creating a copy)
- Use this with caution, as it can lead to unexpected behavior if not handled correctly

■ Memory Management

- Be mindful of memory usage, especially when working with large tensors
- Use techniques like deleting unnecessary tensors and using smaller data types to reduce memory footprint



Layers





What is a Layer ?

- In PyTorch, a "layer" fundamentally refers to a modular building block within a neural network
- It's an object that performs a specific transformation on the input data.
- Think of it like a component in a larger system – each layer has its own purpose and contributes to the overall functionality of the network



What a Layer does ?

- **Linear Transformation:** Applies a matrix multiplication and bias addition (like `nn.Linear`)
- **Convolution:** Performs convolution operations on input data (like `nn.Conv2d`)
- **Pooling:** Downsamples the input data (like `nn.MaxPool2d`)
- **Activation:** Applies a non-linear activation function (e.g., ReLU, Sigmoid). While often used as a separate step in code, activation functions are conceptually part of a layer's transformation
- **Normalization:** Normalizes the input data (e.g., Batch Normalization)
- **Embedding:** Maps discrete values to dense vectors (e.g., `nn.Embedding`)



Characteristics of a Layer

- **Parameters (Learnable):**
 - Most layers have parameters that are learned during training
 - These are typically weights and biases, but can be more complex in some layers
 - These parameters are stored as `torch.Tensor` objects and automatically tracked by PyTorch's optimization process
- **Forward Pass (`forward()` method):**
 - Each layer has a `forward()` method that defines how the input data is transformed into an output
 - This is where the core computation of the layer takes place
- **State (Optional):**
 - Some layers might have internal state that is not learned but influences their behavior
 - For example, RNNs maintain hidden states across time steps
- **Modularity:**
 - Layers are designed to be modular and reusable
 - You can easily combine different layers to create complex network architectures

Types



- **Basic Linear Layers (Fully Connected)**

- **nn.Linear()**: The most fundamental layer which performs a linear transformation

- **Convolutional Layers (For Image/Signal Processing)**

- **nn.Conv2d()**: 2D Convolution. Used for image processing (e.g., CNNs)
- **nn.Conv1d()**: 1D Convolution (for time series data).
- **nn.Conv3d()**: 3D Convolution (for video or volumetric data).

- **Pooling Layers (Downsampling)**

- **nn.MaxPool2d()**: Selects the maximum value within a kernel window
- **nn.AvgPool2d()**: Calculates the average value within a kernel window
- **nn.AdaptiveAvgPool2d()**: Dynamically adjusts the output size to a specified value
- **nn.AdaptiveMaxPool2d()**: Dynamically adjusts the output size to a specified value using max pooling

Types



- **Recurrent Layers (For Sequential Data)**

- **nn.RNN()**: Basic Recurrent Neural Network (RNN) cell
- **nn.LSTM()**: more powerful than basic RNNs for handling long sequences
- **nn.GRU()**: a simplified version of LSTM, often faster to train

- **Embedding Layers**

- **nn.Embedding()**: Maps discrete indices (e.g., word IDs) to dense vectors. Used extensively in natural language processing

- **Normalization Layers**

- **nn.BatchNorm1d()**: Batch Normalization for 1D data
- **nn.BatchNorm2d()**: Batch Normalization for 2D data (e.g., image channels)
- **nn.LayerNorm()**: Layer Normalization – normalizes across features within a layer
- **nn.GroupNorm()**: Group Normalization – normalizes within groups of channels



■ Activation Functions

- **nn.ReLU()**: Rectified Linear Unit – a common activation function
- **nn.Sigmoid()**: Sigmoid function (outputs values between 0 and 1)
- **nn.Tanh()**: Hyperbolic Tangent function (outputs values between -1 and 1)
- **nn.LeakyReLU()**: Leaky ReLU – addresses the "dying ReLU" problem
- **nn.Softmax()**: Softmax function (used in multi-class classification)

■ Dropout Layers (Regularization)

- **nn.Dropout()**: Randomly sets a fraction p of input units to 0 during training.

■ Sequence Operations

- **nn.Flatten()**: Flattens the input tensor into a 1D tensor
- **nn.Unfold()**: Unfolds a 2D or 3D tensor into a sequence
- **nn.Fold()**: Folds a sequence back into a 2D or 3D tensor



Regression





```
import torch

# Define the model
model = torch.nn.Linear(1, 1) # Linear layer: input size 1, output size 1

# Prepare data
X = torch.tensor([[1.0], [2.0], [3.0]]) # Input data
y = torch.tensor([[2.0], [4.0], [6.0]]) # Target data

# Define loss function and optimizer
loss_fn = torch.nn.MSELoss() # Mean Squared Error loss
optimizer = torch.optim.SGD(model.parameters(), lr=0.1) # Stochastic Gradient Descent

# Training loop
for epoch in range(100):
    # Forward pass
    y_pred = model(X)

    # Calculate loss
    loss = loss_fn(y_pred, y)

    # Backward pass and optimization
    optimizer.zero_grad() # Zero the gradients from the previous iteration
    loss.backward() # Calculate gradients
    optimizer.step() # Update parameters

print(f"Loss after training: {loss.item()}")
```



Classification



```
import torch

# Define the model (a simple linear layer)
def model(x, weights, bias):
    """
    Performs a linear transformation followed by a sigmoid activation.

    Args:
        x (torch.Tensor): Input tensor.
        weights (torch.Tensor): Weight matrix.
        bias (torch.Tensor): Bias vector.

    Returns:
        torch.Tensor: Output tensor after transformation and sigmoid activation.
    """
    z = torch.matmul(x, weights) + bias # Linear transformation
    output = torch.sigmoid(z) # Sigmoid activation for classification (0 or 1)
    return output

# Generate synthetic data
num_samples = 100
input_size = 2 # Number of features per sample
output_size = 1 # Binary classification (0 or 1)

X = torch.randn(num_samples, input_size) # Input features
y = (torch.rand(num_samples) > 0.5).float() # Generate random labels (0 or 1)

# Initialize weights and bias randomly
weights = torch.randn(input_size, output_size)
bias = torch.randn(output_size)

# Define hyperparameters
learning_rate = 0.1
epochs = 100
```




```
# Training loop using PyTorch's loss function and optimizer
criterion = torch.nn.functional.binary_cross_entropy # Binary Cross Entropy Loss
optimizer = torch.optim.SGD(params=[weights, bias], lr=learning_rate) # Stochastic Gradient Descent

for epoch in range(epochs):
    # Forward pass
    predictions = model(X, weights, bias)

    # Calculate loss
    loss = criterion(predictions, y)

    # Backpropagation and optimization
    optimizer.zero_grad() # Zero the gradients before each backward pass
    loss.backward() # Calculate gradients
    optimizer.step() # Update weights and bias

    # Print loss every 10 epochs
    if (epoch + 1) % 10 == 0:
        print(f"Epoch {epoch+1}, Loss: {loss.item()}")

# Evaluation (after training)
with torch.no_grad(): # Disable gradient calculation during evaluation
    predictions = model(X, weights, bias)
    predicted_labels = (predictions > 0.5).float() # Convert probabilities to labels
    accuracy = torch.sum(predicted_labels == y) / num_samples

print(f"Accuracy: {accuracy.item()}")
```